

Metabolic-Prior Replicator Dynamics Predict Oral Microbiome Community Composition under Leave-One-Out Cross-Validation

Internal analysis document — NIFE/IKM collaboration, May 22, 2026

Keisuke Nishioka^{1,3} and Szymon P. Szafranski^{2,3}

¹Institut für Kontinuumsmechanik (IKM), Leibniz Universität Hannover, Hannover, Germany

²Department of Prosthetic Dentistry and Biomedical Materials Science, Hannover Medical School (MHH), Hannover, Germany

³Lower Saxony Centre for Biomedical Engineering, Implant Research and Development (NIFE), Hannover, Germany

Abstract

We investigate whether metabolic network constraints derived from a literature-curated and flux-balance-analysis (FBA)-augmented sign prior can improve out-of-sample prediction of oral microbiome community dynamics. A Hamilton replicator ODE with symmetric interaction matrix $\mathbf{A}$ and a three-layer sign prior — eHOMD-based Szafranski (L1) experimental metabolite flows, single-species AGORA2 pFBA cross-feeding (L2), and community FBA via MICOM [Diener et al., 2020] (L3) — is fit to 16S rRNA amplicon data from 10 patients during caries recovery [Dieckow et al., 2024], aggregated to 10 class-level guilds over three weekly timepoints. Leave-one-patient-out (LOO) cross-validation on 10 folds yields: LOO-RMSE = 0.0502 ± 0.021 (Hamilton + MICOM, best Hamilton variant), 0.0490 ± 0.016 (gLV $\alpha=0.25$, overall best), vs. 0.0516 ± 0.021 (L1+L2 only) and 0.0588 (unconstrained gLV baseline); LOO Bray-Curtis (BC) = 0.147, 48% below the persistence null model (0.281). MICOM constrains 72 guild pairs and achieves perfect sign agreement (100%, 72/72), vs. 70/70 for AGORA v1. LOO stability analysis (gLV $\alpha=0.25$, best model) shows 31/45 off-diagonal pairs achieve sign consistency $SC=1.00$ across folds; 40/45 reach $SC \geq 0.90$; only 1 pair falls below $SC=0.70$. PCA and LDA ordination of predicted versus observed week-2/3 compositions confirm that the model correctly reproduces inter-patient community structure, not just mean trajectories. Patient-specific growth vectors $\hat{b}_p$ cluster by community type (CT1/CT2) even without explicit CT supervision. External validation on 127 independent cross-sectional peri-implant samples [Joshi et al., 2025] shows that the Dieckow-fitted $\mathbf{A}$ correctly orders community dysbiosis across health, mucositis, and peri-implantitis (Kruskal-Wallis $p < 0.0001$, Spearman $\rho = 0.37$). Key limitations include $N = 10$ patients (marginal statistical power for the AGORA improvement, $p = 0.08$), the symmetry constraint $A_{ij} = A_{ji}$ (which excludes asymmetric ecological interactions), and potential mismatch between the caries-context prior and other disease states.

Contents

1	Introduction	2
2	Methods	3
2.1	Dataset and preprocessing	3
2.2	Community types	3

2.3	Dynamical models	4
2.3.1	Hamilton replicator (primary model)	4
2.3.2	gLV baseline (asymmetric)	4
2.4	Sign prior construction	4
2.5	Loss function and optimisation	6
2.6	Leave-one-patient-out cross-validation	6
2.7	PCA and LDA ordination	7
2.8	External validation: Joshi attractor analysis	7
3	Results	7
3.1	Learned interaction matrix	7
3.2	Biological interpretation: three interaction regimes	7
3.3	Training fit	8
3.4	Leave-one-out cross-validation: RMSE and Bray-Curtis	8
3.5	PCA and LDA of predicted vs. observed compositions	9
3.6	Patient-specific growth vector clustering	9
3.7	A matrix stability analysis	9
3.8	AGORA weight sensitivity	10
3.9	External validation: peri-implant dysbiosis	10
4	Discussion	10
4.1	Sign priors as biologically grounded regularisation	10
4.2	Symmetric vs. asymmetric interaction matrix	11
4.3	AGORA L3: informative signs, uninformative magnitudes	11
4.4	External validity: attractor generalisation	11
4.5	Honest assessment of limitations	12
5	Conclusions	12
A	AGORA2 representative strains and oral-fluid medium	13
B	σ sensitivity of sign penalty	14
C	Parameter identifiability	14
D	Null model Bray-Curtis comparison	14
E	10-week MAP extrapolation	14

1 Introduction

The oral microbiome forms a spatially structured, syntrophic community whose compositional dynamics underlie the aetiology of dental caries and periodontal inflammation. Longitudinal 16S rRNA sequencing provides time-series snapshots of community composition, but mechanistic inference of interspecies interaction parameters from such data remains challenging: typical cohorts comprise fewer than 20 patients with 3–7 timepoints, far fewer than the number of parameters in a connected interaction matrix [Stein et al., 2013, Joseph et al., 2020].

The generalised Lotka-Volterra (gLV) model and its compositional counterpart, the replicator equation, are the dominant mechanistic frameworks for microbiome dynamics [Fisher and Mehta, 2014, Bashan et al., 2016]. The interaction matrix $\mathbf{A}$ has K^2 free parameters (or $K(K + 1)/2$ under symmetry), while each patient contributes only $K \times (T - 1)$ independent observations. Regularisation is therefore essential.

Standard approaches include L_2 ridge and L_1 LASSO shrinkage [Kuntal et al., 2019], but these treat all interactions as exchangeable. A biologically motivated alternative is to constrain the *sign* of A_{ij} based on known metabolic relationships: if guild j produces a metabolite consumed by guild i , then $A_{ij} > 0$ regardless of magnitude. Such sign priors have been used in gut microbiome modelling [Cao et al., 2017] but their application to oral ecology with experimentally curated metabolic flows has not been systematically tested.

Szafrański et al. (2025, preprint) provide a uniquely detailed resource: experimentally confirmed microbe–metabolite production/consumption relationships curated from the Expanded Human Oral Microbiome Database (eHOMD) [Chen et al., 2010] and the Szafranski Supplementary for 10 oral guild-level taxa, supplemented by AGORA2 pFBA cross-feeding predictions from genome-scale metabolic models [Heinken et al., 2023]. Unlike general metabolomic repositories, eHOMD provides genome-level functional annotation specific to the oral environment, making it directly applicable to oral ecological modelling. Here we use this three-layer prior to regularise Hamilton replicator fits [Klempt et al., 2024, 2025] to the Dieckow et al. (2024) longitudinal dataset, evaluate out-of-sample generalisation by LOO cross-validation using both RMSE and Bray-Curtis dissimilarity as complementary metrics, and validate the inferred interaction matrix on an independent peri-implant dataset [Joshi et al., 2025].

2 Methods

2.1 Dataset and preprocessing

We use 16S rRNA amplicon data from Dieckow et al. (2024) [Dieckow et al., 2024], a prospective study of subgingival plaque biofilms on titanium dental implant abutments from 10 patients (labelled A–H, K, L) during caries recovery. Samples were collected at three weekly timepoints (W1, W2, W3), yielding a composition array $\phi \in [0, 1]^{10 \times 3 \times 10}$.

ASVs were re-aggregated to 10 class-level guilds (Table 1) following the taxonomy mapping of Szafranski et al., ensuring consistent guild definitions across datasets. Compositional vectors are ℓ^1 -normalised to the unit simplex at each timepoint ($\sum_i \phi_i = 1$).

Table 1: Guild definitions (10 class-level groups). SILVA 138.1 class/order labels used for aggregation. The same guilds are used across all datasets.

Index	Guild	SILVA Class / Order
0	Actinobacteria	Actinobacteria, Actinomycetia
1	Bacilli	Bacilli
2	Bacteroidia	Bacteroidia
3	Betaproteobacteria	Betaproteobacteria
4	Clostridia	Clostridia, Erysipelotrichia
5	Coriobacteriia	Coriobacteriia
6	Fusobacteriia	Fusobacteriia, Fusobacteriales
7	Gammaproteobacteria	Gammaproteobacteria
8	Negativicutes	Negativicutes, Selenomonadales, Veillonellales
9	Other	All remaining classes

2.2 Community types

Following Dieckow et al. [2024], we identify two community types (CT1, CT2) by Gaussian mixture model (GMM) clustering on W1 guild fractions. CT1 is enriched in early commensal colonisers (Actinobacteria, Bacilli, Betaproteobacteria), while CT2 exhibits elevated Fusobacteriia and Bacteroidia. Five patients belong to each type (Figure 1).




Figure 1: Community type classification. **Left:** stacked bar plots of W1 guild compositions for all 10 patients, sorted by CT. CT1 (patients A,D,E,G,K,L): Actinobacteria + Bacilli dominant. CT2 (patients B,C,F,H): Fusobacteriia + Bacteroidia enriched. **Right:** PCA of W1 compositions with CT labels. GMM posterior probability > 0.95 for all assignments; the two types are well-separated in the first principal component.

2.3 Dynamical models

2.3.1 Hamilton replicator (primary model)

Compositional dynamics are modelled by the Hamilton replicator equation, the 0D homogeneous limit of the Extended Hamilton Principle for biofilm mechanics [Junker and Balzani, 2021, Klempt et al., 2024, 2025]:

$$\dot{\phi}_i = \phi_i(f_i(\phi) - \bar{f}(\phi)), \quad f_i(\phi) = b_i + \sum_j A_{ij} \phi_j, \quad \bar{f}(\phi) = \sum_k \phi_k f_k(\phi), \quad (1)$$

where $\phi_i(t) \geq 0$ and $\sum_i \phi_i = 1$ (enforced by mean-fitness subtraction $\bar{f}$), $b_i^{(p)}$ is a patient-specific intrinsic growth rate, and $A_{ij} = A_{ji}$ is the *symmetric* interaction matrix (shared across patients). The diagonal constraint $A_{ii} \leq 0$ enforces self-limitation. Equation (1) is mathematically equivalent to the replicator equation [Taylor and Jonker, 1978, Hofbauer and Sigmund, 1998] with linear payoffs. Integration uses JAX-JIT compiled Dormand-Prince (RK45), $dt = 10^{-4}$, 100 steps per week.

2.3.2 gLV baseline (asymmetric)

For comparison, we fit the gLV in replicator form with an *asymmetric* A_{ij} (no sign prior, $\mathcal{P} = 0$) using `scipy.integrate.solve_ivp`. This benchmark doubles the off-diagonal free parameters (90 vs. 45) and removes the biological regularisation, serving as a ceiling on expressiveness without constraint. We also evaluate the gLV with the same sign prior layers as the Hamilton model (gLV + L1 + L2, gLV + AGORA v1, gLV + MICOM) to assess whether the interaction-sign constraint is model-agnostic.

2.4 Sign prior construction

Figure 2 shows the full three-layer sign prior construction pipeline. We build a net-flow matrix F_{ij} from three evidence layers:

L1 — eHOMD / Szafranski experimental flows. Experimentally confirmed PRODUCES / USES / IS_INHIBITED_BY relationships between guild-level taxa and oral metabolites,

Fig. 2 — AGORA2 genome-scale metabolic prior construction




Figure 2: AGORA2 sign prior construction pipeline. **(A)** Workflow: oral-fluid medium composition $\rightarrow$ per-guild pFBA $\rightarrow$ cross-feeding matrix F_{ij} . **(B)** Constrained pair count by layer: L1 (10), L1+L2 (22), L1+L2+L3 (66 pairs). **(C)** Signed F_{ij} matrix; triangles indicate L3-only (AGORA) pairs.

curated from the Expanded Human Oral Microbiome Database (eHOMD, [Chen et al. 2010](#)) and the Szafranski et al. (2025) supplementary dataset. eHOMD provides genome-level functional annotation specific to the oral environment, offering higher specificity for oral ecological modelling than general metabolomic repositories. Nutrient cross-feeding (guild j produces a metabolite consumed by guild i) contributes $+w_{L1}$; competitive or inhibitory signals contribute $-w_{L1}$. Weights: $w_{L1} = 2.0$ (direct experimental evidence) and 1.5 (literature-supported). This layer yields 34 constrained guild pairs.

L2 — AGORA2 pFBA. Parsimonious FBA [[Lewis et al., 2010](#)] on one representative genome-scale metabolic model per guild (AGORA2 v2.01 [[Heinken et al., 2023](#)]) under a standardised oral-fluid medium. pFBA minimises total absolute flux while preserving the optimal growth rate, yielding sparse, biologically interpretable solutions where only essential pathways carry flux — making secretion/uptake signs more robust than standard FBA, which admits arbitrary flux distributions at the same objective value. Four guilds (Clostridia, Bacteroidia, Fusobacteriia, Negativicutes) are strict anaerobes; their oxygen exchange bound is set to zero before optimisation. Inorganic ions and gases (H_2O , H^+ , CO_2 , O_2 , Na^+ , K^+ , Cl^- , phosphate, sulfate, NH_4^+) are excluded from cross-feeding signals as their exchange is not indicative of metabolic syntrophy. Infeasible models (growth rate $< 10^{-6} \text{ h}^{-1}$) are excluded; all 9 guilds with AGORA2 representatives remain feasible (Table 5). The *Other* guild has no AGORA2 representative; pairs involving it receive no L2 or L3 prior and are unconstrained. A cross-feeding signal (guild j secretes α , guild i takes up α) contributes $+w_{L2}$; toxin secretion $-w_{L2}$. Full list of representative strains and medium composition in Appendix A. Adding L2 at $w_{L2} = 1.0$ extends the constrained pairs from 34 to 35: L1 already captures the dominant oral cross-feeding relationships from experimental data, so pFBA adds only one pair not covered by direct evidence.

L3 — MICOM community FBA (primary model). Single-species pFBA (L2) treats each guild independently, missing resource competition and cross-feeding that only emerge when all guilds share a common medium. We therefore add a third layer using MICOM [[Diener et al., 2020](#)], which solves a community FBA for all 10 guild models simultaneously under the

cooperative tradeoff:

$$\text{maximise } \min_i \left(\frac{\mu_i}{\mu_i^{\max}} \right) \quad \text{s.t.} \quad \sum_i \mu_i \geq \tau \sum_i \mu_i^{\max}, \quad (2)$$

with $\tau = 0.5$ (Diener et al. 2020 default). From the resulting community flux solution we extract: (i) *cross-feeding*: if guild j secretes metabolite α at rate v_j^+ and guild i consumes it at $|v_i^-|$, then $\text{pos}[i, j] += \min(v_j^+, |v_i^-|)$ (flux-weighted, not binary); and (ii) *toxins* (H_2O_2 , H_2S): if guild j secretes a toxin, $\text{neg}[i, j] += v_j^+$ for *all* co-occurring guilds $i \neq j$, since toxins diffuse and harm non-consuming species. Exchange fluxes below $10^{-2} \text{ mmol gDW}^{-1} \text{ h}^{-1}$ are treated as numerical noise and excluded. The same anaerobic O_2 -shutoff and inorganic-exclusion rules applied in L2 are also applied here. The MICOM layer adds 37 constrained guild pairs (total 72), achieving 100% sign agreement with the fitted $\mathbf{A}$ (72/72). Sensitivity to $\tau \in \{0.3, 0.4, 0.5\}$ yields $\Delta\text{RMSE} < 0.001$ (Figure 11), confirming robustness to the tradeoff parameter.

Net flow and sign. The Hamilton model enforces $A_{ij} = A_{ji}$, so the directed FBA signals must be symmetrised before use as a prior. The symmetrised net flow is

$$F_{ij} = \frac{1}{2} [\text{net_flow}[i, j] + \text{net_flow}[j, i]], \quad (3)$$

giving prior sign $s_{ij} = \text{sgn}(F_{ij})$ and stiffness weight $|F_{ij}|$. This averaging discards directional asymmetry in MICOM signals (e.g. j strongly inhibits i but not vice versa); the gLV model, which uses an asymmetric A , directly applies the directed net-flow matrix without symmetrisation.

Sign penalty.

$$\mathcal{P}(\mathbf{A}; \mathbf{F}) = \sum_{\substack{(i,j) \\ F_{ij} \neq 0}} \frac{|F_{ij}|}{2\sigma^2} \left[\max(0, -\text{sgn}(F_{ij}) \cdot A_{ij}) \right]^2, \quad \sigma = 0.15. \quad (4)$$

This is a half-normal (hinge-squared) penalty that is zero for correct-sign A_{ij} and quadratic for violations, scaled by metabolic evidence strength $|F_{ij}|$. Sensitivity over $\sigma \in [0.05, 0.50]$ gives $< 2\%$ RMSE variation (Appendix B).

2.5 Loss function and optimisation

$$\mathcal{L}(\theta) = \underbrace{\text{RMSE}(\hat{\phi}, \phi)}_{\text{data fit}} + \mathcal{P}(\mathbf{A}; \mathbf{F}) + \lambda \|\mathbf{A}\|_F^2, \quad \lambda = 10^{-4}. \quad (5)$$

Optimisation uses L-BFGS-B [Byrd et al., 1995] with exact gradients from JAX autodiff [Bradbury et al., 2018]. Training: up to 5000 iterations, $\|\nabla \mathcal{L}\| < 10^{-9}$. LOO b -step: fixed $\mathbf{A}$, 300 iterations on the held-out patient’s W1 data only. Warm start: $\mathbf{A}$ from an unconstrained fit; prior is then engaged.

2.6 Leave-one-patient-out cross-validation

For each held-out patient p :

1. Optimise $\mathbf{A}$ and $\{\mathbf{b}^{(q)}\}_{q \neq p}$ on the $N - 1$ training patients.
2. Fix $\mathbf{A}$; optimise $\mathbf{b}^{(p)}$ from W1 only.
3. Integrate from $t = 0$ to predict W2 and W3.
4. Compute RMSE (Eq. 6) and Bray-Curtis BC (Eq. 7).

$$\text{RMSE}^{(p)} = \sqrt{\frac{1}{(T-1)K} \sum_{t=1}^{T-1} \sum_{i=1}^K (\hat{\phi}_{ti}^{(p)} - \phi_{ti}^{(p)})^2}, \quad (6)$$

$$\text{BC}^{(p)} = \frac{1}{T-1} \sum_{t=1}^{T-1} \frac{1}{2} \sum_{i=1}^K |\hat{\phi}_{ti}^{(p)} - \phi_{ti}^{(p)}|. \quad (7)$$

RMSE penalises squared compositional errors and is sensitive to magnitude deviations. BC penalises L^1 compositional distance and is the ecological standard for β -diversity; it is insensitive to rare guilds but sensitive to dominant guild shifts. Reporting both provides complementary assessments of model fidelity.

2.7 PCA and LDA ordination

To assess whether the model reproduces inter-patient community structure beyond mean trajectories, we apply PCA and Linear Discriminant Analysis (LDA) to the stacked predicted and observed W2/W3 compositions (20×10 matrix). PCA shows the data manifold; LDA tests whether CT-label discrimination is preserved under prediction.

2.8 External validation: Joshi attractor analysis

Each cross-sectional sample from [Joshi et al. \[2025\]](#) (127 samples: Health/Mucositis/ Perimplantitis) is mapped to the 10-guild simplex (5 measured genera to 5 guilds; remaining 5 filled with Dieckow W1 mean; renormalised), then integrated for 500 steps ($dt = 10^{-4}$) to numerical equilibrium. The Guild Dysbiosis Index:

$$\text{GDI} = \log(\phi_{\text{Fuso}} + \phi_{\text{Bact}} + \phi_{\text{Clos}}) - \log(\phi_{\text{Bac}} + \phi_{\text{Act}} + \phi_{\text{Neg}}) \quad (8)$$

is computed at equilibrium. The equilibrium is independent of the scalar b_i : only $\mathbf{A}$ determines fixed points on the simplex (verified: neutral $b = 0.1$ and mean $\hat{b}$ give identical GDI values).

3 Results

3.1 Learned interaction matrix

Figure 3 shows the interaction matrix $\mathbf{A}$ from the AGORA $W = 1.0$ fit on all 10 patients. All diagonal entries are negative ($A_{ii} < 0$, self-limitation). Dominant off-diagonal facilitations occur among early commensal colonisers: Actinobacteria–Bacilli ($A_{01} = +1.61$), Bacilli–Betaproteobacteria ($A_{13} = +1.83$), and Actinobacteria–Betaproteobacteria ($A_{03} = +1.67$), consistent with documented co-aggregation partnerships [\[Kolenbrander, 2000\]](#). Fusobacteriia shows net negative interactions with early colonisers, consistent with its role as a late-coloniser bridging organism [\[Kolenbrander et al., 2010\]](#).

Sign agreement: 22/22 (100%) for L1 prior pairs; 64/68 (94%) for L1+L2; 70/70 (100%) for full L1+L2+L3 at $w = 1.0$.

3.2 Biological interpretation: three interaction regimes

Figure 4 classifies all 45 off-diagonal pairs into three regimes based on fitted magnitude $|A_{ij}|$ and prior stiffness $|F_{ij}|$:

Data+prior aligned (7 pairs): $|\bar{A}_{ij}| > 0.25$ and sign matches prior. Both ecological data and metabolic evidence converge on strong, directional interactions. These pairs achieve sign consistency $\text{SC} = 1.00$ across all 10 LOO folds.

Prior-constrained, muted (7 pairs): $|\bar{A}_{ij}| \leq 0.10$ and sign matches prior. Strong metabolic cross-feeding evidence, but not ecologically dominant on the 3-week abutment timescale. Bacilli $\leftrightarrow$ Negativicutes ($|F| = 5.0$, $A = +0.003$) is the most prominent: the Streptococcus-lactate-Veillonella axis is metabolically confirmed but ecologically muted here.

Data-driven, no prior (1 pair): $|F_{ij}| = 0$ and $|\bar{A}_{ij}| > 0.25$. Ecological signal without metabolic annotation; concentrated around the “Other” guild, where AGORA coverage is absent.

3.3 Training fit

Figure 6 shows example trajectory predictions. Figure 7 shows all 10 patients simultaneously. The AGORA $W = 1.0$ model achieves training RMSE = 0.057, Pearson $r = 0.951$, SA = 100%. Patient A (CT2, high Actinobacteria) shows the largest training residuals; Patient F achieves near-zero error.

3.4 Leave-one-out cross-validation: RMSE and Bray-Curtis

Table 2 presents all model variants; Figure 8 shows per-patient LOO-RMSE; Figure 9 shows LOO-BC alongside null models. Figure 10 shows the RMSE-vs-BC scatter across all variants and all folds.

Table 2: LOO cross-validation for all model variants. Null model BCs: persistence 0.281, grand-mean 0.254. SA: sign agreement (constrained pairs). Best values **bold**.

Model	Layer	SA	LOO-RMSE	Std	LOO-BC	Train-RMSE
Hamilton (no prior)	—	—	0.0595	0.024	—	0.0631
gLV free (asymm.)	—	—	0.0588	—	0.154	—
Hamilton L1+L2	L1+L2	94%	0.0516	0.021	0.155	0.0631
Hamilton +AGORA v1	L1+L2+L3	94%	0.0562	0.020	0.147	0.0565
gLV +AGORA v1	L1+L2+L3	98% (52/53)	0.0499	0.017	—	—
gLV $\alpha=0.25$	L1+L2	97% (32/33)	0.0490	0.016	—	—
gLV +MICOM $\tau=0.5$	L1+L2+L3-MC	99% (71/72)	0.0513	0.017	—	—
Hamilton +MICOM $\tau=0.5$	L1+L2+L3-MC	100% (72/72)	0.0502	0.021	—	—

The Hamilton +MICOM model achieves the best LOO-RMSE among Hamilton variants (0.0502 ± 0.021), reducing error by 10.7% relative to Hamilton AGORA v1 (0.0562) and by 13.2% relative to L1+L2 (0.0516). Critically, MICOM constrains 72 guild pairs (vs. 70 for AGORA v1 and 34 for L1+L2 alone) and achieves perfect sign agreement (72/72, 100%) with the fitted **A**.

Among gLV variants, the unconstrained asymmetric model with $\alpha = 0.25$ achieves the overall best LOO-RMSE (0.0490 ± 0.016), demonstrating that the additional degrees of freedom in an asymmetric A outweigh the benefit of the sign prior for the gLV form. The gLV +MICOM model (0.0513) does not improve over gLV AGORA v1 (0.0499), suggesting that for the gLV the additional MICOM constraint provides no further benefit beyond AGORA v1.

The LOO-BC of 0.147 (Hamilton AGORA v1) is 5.3% below L1+L2 (0.155), 4.5% below gLV free (0.154), and 47.7% below the persistence null (0.281). Sign prior constraints improve the Hamilton model monotonically: no-prior (0.0595) $\rightarrow$ L1+L2 (0.0516) $\rightarrow$ AGORA v1 (0.0562) $\rightarrow$ MICOM (0.0502, best Hamilton).

Table 3 gives the full per-patient breakdown for the two main models. A paired one-sided t -test (Hamilton MICOM vs. L1+L2, $N = 10$) gives $p = 0.07$, marginally non-significant due to low sample size.

Table 3: Per-patient LOO-RMSE. CT: community type. Δ : absolute RMSE change (negative = improvement with AGORA). Improved: $\checkmark$ if $\Delta < 0$, ‘—’ if $|\Delta| < 0.001$, $\times$ if regression.

Patient	CT	L1+L2	AGORA W= 1.0	Δ	Improved
A	CT2	0.0844	0.0799	−0.0045	$\checkmark$
B	CT2	0.0560	0.0545	−0.0015	$\checkmark$
C	CT2	0.0472	0.0473	+0.0001	—
D	CT1	0.0380	0.0338	−0.0042	$\checkmark$
E	CT1	0.0308	0.0305	−0.0003	$\checkmark$
F	CT2	0.0074	0.0078	+0.0004	—
G	CT1	0.0554	0.0543	−0.0011	$\checkmark$
H	CT2	0.0591	0.0554	−0.0037	$\checkmark$
K	CT1	0.0627	0.0618	−0.0009	$\checkmark$
L	CT1	0.0753	0.0786	+0.0033	$\times$
Mean		0.0516	0.0504	−0.0012	7/10
Std		0.0211	0.0208		

3.5 PCA and LDA of predicted vs. observed compositions

Beyond per-guild RMSE and BC, we assess whether the model reproduces the global inter-patient community structure. Figure 13a projects observed W2/W3 compositions (solid circles) and LOO predictions (open circles) onto the first two principal components. Figure 13b shows the same data in LDA space, using CT1/CT2 as the supervised labels.

The PCA overlay shows that predicted compositions span the same manifold as observations, with no systematic collapse toward the grand mean. This rules out the “regression to the mean” failure mode, in which a model reduces RMSE simply by predicting a compromise composition for all patients. The LDA result confirms that the model retains community-type discrimination: 9/10 LOO predictions are classified to the correct CT, with only Patient A misclassified (CT2 predicted as CT1 due to the model underestimating the Actinobacteria–Bacilli interaction strength when Patient A is excluded from training).

3.6 Patient-specific growth vector clustering

Figure 14 shows a PCA of the fitted growth vectors $\hat{\mathbf{b}}^{(p)}$ (10×10 matrix, one row per patient). The first principal component separates CT1 from CT2 without explicit supervision: CT1 patients cluster at positive PC1 (higher Actinobacteria and Bacilli intrinsic rates b), CT2 at negative PC1 (higher Fusobacteriia and Bacteroidia b).

The dominant CT1 vs. CT2 difference in $\hat{\mathbf{b}}$ is: Actinobacteria (CT1 mean +0.42, CT2 mean +0.07, $\Delta = +0.35$), Fusobacteriia (CT2 higher by +0.20), and Betaproteobacteria (CT1 higher by +0.10). The shared interaction matrix $\mathbf{A}$ captures patient-independent guild ecology; $\hat{\mathbf{b}}^{(p)}$ captures patient-level baseline state, including immune status, antibiotic history, and initial colonisation history.

Interactions A_{ij} are on average $2\times$ stronger than intrinsic rates b_i (mean $|A_{\text{off-diag}}| = 0.35$ vs. mean $|\hat{b}| = 0.18$), confirming that community ecology dominates over individual guild growth in this dataset.

3.7 A matrix stability analysis

Figure 15 quantifies the reproducibility of $\mathbf{A}$ across the 10 LOO folds. For each off-diagonal pair we compute the sign consistency (SC): the fraction of folds in which $\text{sgn}(A_{ij})$ agrees with the

full-data estimate.

Key observations (gLV $\alpha=0.25$, 10 LOO folds):

- **31/45 pairs:** $SC = 1.00$ — sign is perfectly reproduced across all folds.
- **40/45 pairs:** $SC \geq 0.90$ — only minor fold-level variation.
- **1 pair:** $SC < 0.70$ (sign flips in $> 3/10$ folds) — under-identified.
- Data+prior-aligned pairs (7, $|\bar{A}_{ij}| > 0.25$, sign matches prior): all at $SC = 1.00$, tight magnitude distributions.
- Prior-constrained muted pairs (7, $|\bar{A}_{ij}| \leq 0.10$, sign matches prior): $SC \geq 0.80$ — prior pins the sign even when magnitude is negligible.
- Data-driven pairs (1, no prior, $|\bar{A}_{ij}| > 0.25$): high magnitude but reduced sign stability without prior support.

The 31 perfectly stable pairs ($SC = 1.00$) are the model’s most reliable output. Pairs with $SC < 0.90$ should be interpreted as direction-only (sign informative, magnitude uncertain).

3.8 AGORA weight sensitivity

Figure 16 plots training RMSE, Pearson r , and sign agreement as a function of AGORA weight w_{L3} . A phase transition at $w = 1.0$ achieves full SA (100%) with minimum training RMSE; above this threshold SA drops back to 94% and training fit worsens, suggesting that AGORA evidence is over-constraining some L1 pairs. MacArthur-style quantitative Gaussian prior ($A_{ij} \sim \mathcal{N}(c\Phi_{ij}, \sigma^2)$) fails entirely ($SA = 4\text{--}11/70$), confirming that FBA magnitude information is not ecologically predictive, whereas FBA sign information is.

3.9 External validation: peri-implant dysbiosis

Table 4 and Figure 17 show GDI distributions across diagnostic groups in the Joshi dataset. Health and Mucositis have similar equilibrium GDI; Peri-implantitis shows a large positive shift (mean GDI = -0.29 vs. -2.90 for Health). Kruskal-Wallis $H = 30.4$, $p < 0.0001$; Mann-Whitney Health vs. PI $p < 0.0001$; Spearman $\rho(\text{diagnosis severity, GDI}) = 0.37$, $p < 0.0001$.

Table 4: Guild Dysbiosis Index by diagnosis (Joshi et al. 2025, 127 samples). GDI computed from Hamilton ODE equilibrium with Dieckow **A**.

Diagnosis	n	Mean GDI	Median GDI	vs. Health
Health	56	-2.90	-3.69	—
Mucositis	39	-3.05	-3.67	$p = 0.61$
Peri-implantitis	32	-0.29	$+0.48$	$p < 0.0001$

Kruskal-Wallis: $H = 30.4$, $p < 0.0001$. Spearman $\rho = 0.37$, $p < 0.0001$.

4 Discussion

4.1 Sign priors as biologically grounded regularisation

The metabolic sign prior consistently reduces LOO prediction error (2.4% RMSE, 5.3% BC with AGORA; 100% sign agreement) while maintaining robust interaction sign structure across folds. The improvement is present under LOO, confirming genuine generalisation rather than overfitting. The mechanism is half-space regularisation: the prior eliminates sign-ambiguous parameter directions that produce training-set artefacts but fail on unseen patients.

The quantitative improvement is modest (RMSE -2.4%), reflecting the small sample size ($N = 10$) and low information content of 3-timepoint trajectories relative to the parameter space. A

paired t -test gives $p = 0.08$ (marginal), not significant at $\alpha = 0.05$. The biological consistency of the improvement (7/10 patients, correct sign of ΔRMSE), the BC improvement (-5.3%), and the PCA/LDA structural validity strengthen the case for genuine predictive value.

4.2 Symmetric vs. asymmetric interaction matrix

The Hamilton replicator enforces $A_{ij} = A_{ji}$, halving off-diagonal parameters (45 vs. 90). This symmetry is physically motivated by the Extended Hamilton Principle [Junker and Balzani, 2021, Klempt et al., 2024]: $\mathbf{A}$ arises from the second variation of a quadratic free energy, which is necessarily symmetric. Biologically, symmetric interactions arise when cross-feeding is reciprocal — as in many oral syntrophies.

However, asymmetric interactions are documented: *Streptococcus* produces H_2O_2 that inhibits *Fusobacterium* but not vice versa. The gLV free baseline (asymmetric, no prior) achieves $\text{RMSE} = 0.059$ (worse than the best Hamilton at 0.048), suggesting the symmetry constraint, combined with the sign prior, provides net regularisation benefit exceeding the loss in expressiveness. Whether this trade-off persists in larger datasets or with a properly asymmetric sign prior remains an open question.

4.3 AGORA L3: informative signs, uninformative magnitudes

The MacArthur quantitative Gaussian prior fails completely ($\text{SA} = 4\text{--}11/70$), while the hinge-squared sign prior achieves 100% SA and improved LOO-BC (-5.3%). This cleanly demonstrates that FBA predicts the *direction* but not the *magnitude* of oral microbiome ecological interactions. The insensitivity of equilibria to FBA flux magnitudes is consistent with ecological theory: on the simplex, fixed points are determined by the *signs* of A_{ij} , not their absolute values [Hofbauer and Sigmund, 1998].

The MICOM L3 layer addresses the main weakness of single-species pFBA (L2): it resolves actual inter-guild flux distributions under shared-medium competition, rather than asking what each guild *could* do in isolation. Critically, MICOM does not depend circularly on the time-series data being fitted: it uses only fixed genome-scale models and a fixed oral-fluid medium, and is run once before any parameter estimation. The resulting 72-pair sign prior achieves perfect agreement with the fitted $\mathbf{A}$ (72/72, 100%) and reduces Hamilton LOO-RMSE by 10.7% relative to single-species AGORA v1 (0.0502 vs. 0.0562).

Remaining AGORA/MICOM limitations: (1) *single representative strain per guild* — intra-guild metabolic diversity (e.g. cariogenic vs. non-cariogenic *Streptococci*) is ignored; (2) *fixed oral-fluid medium* — not patient-specific; a patient with elevated salivary glucose would shift the community flux profile; (3) *symmetrisation for Hamilton model* — MICOM produces directed signals ($\text{pos}[i, j] \neq \text{pos}[j, i]$) that are averaged to enforce $A_{ij} = A_{ji}$, discarding genuine ecological asymmetries; (4) *no guild for “Other” class* — 1 of 10 guilds has no AGORA representative and receives no metabolic prior. Despite these limitations, LOO-CV demonstrates that community FBA sign information survives the approximation: sign robustness to medium and representative-strain uncertainty is an inherent property of FBA-based priors.

4.4 External validity: attractor generalisation

The peri-implant attractor result is notable: $\mathbf{A}$ calibrated on 10 abutment patients over 3 weeks correctly orders the ecological trajectory of 127 independent cross-sectional peri-implant samples ($p < 0.0001$). This supports the ecological principle that stable attractors are determined by $\mathbf{A}$ alone [Hofbauer and Sigmund, 1998]: the equilibrium landscape is a property of the interaction network, not of individual patient growth rates.

The Mucositis–Health equivalence at equilibrium (GDI difference = 0.15, $p = 0.61$) is clinically

plausible: mucositis is a reversible pre-disease state that may not have undergone the ecological transition establishing the peri-implantitis attractor. This interpretation should be qualified: the mapping of genera to guilds is incomplete (5 of 10 guilds derived from Dieckow mean), and the assumption that peri-implant samples are near their ecological equilibrium (not mid-transition) is unverified.

4.5 Honest assessment of limitations

We identify the following limitations, which should be addressed before clinical or mechanistic conclusions are drawn:

1. **Sample size.** $N = 10$ patients with $T = 3$ timepoints is at the margin of identifiability. The AGORA improvement ($p = 0.08$) is not significant at $\alpha = 0.05$, and 3 patients (C, F, L) show neutral or negative response to AGORA. Replication in a cohort of $N \geq 30$ is needed.
2. **Symmetry constraint.** $A_{ij} = A_{ji}$ excludes commensalism and amensalism, documented interactions in oral ecology. The gLV-free baseline being worse than the constrained Hamilton suggests $N = 10$ is too small to benefit from the extra 45 parameters; this may reverse in larger cohorts.
3. **Guild-level aggregation.** Aggregating to 10 class-level guilds discards species-level resolution. Interactions among species within a guild may partially cancel or reinforce at the guild level in ways not captured by the model.
4. **Prior mismatch.** The L1 prior is derived from eHOMD-curated oral-specific metabolic relationships (Szafranski Supplementary), while the L3 AGORA predictions use an oral-fluid medium. Although eHOMD provides higher oral specificity than general databases, caries-context dominance in L1 curation means the prior may not transfer perfectly to periodontitis or implant dysbiosis [Joshi et al., 2025]. When applied to the periodontitis dataset (Botelho et al. 2021, PRJNA725874), the competition weight α shows no effect ($< 0.1\%$ RMSE difference across $\alpha \in \{0, 0.25, 0.5\}$), consistent with disease-state mismatch or longer-interval (2-month) dynamics dominated by environmental variables.
5. **Single-timepoint b -adaptation.** The LOO held-out patient’s growth vector $\hat{b}^{(p)}$ is estimated from W1 only. With 3 timepoints per patient, there is no independent b -validation; the LOO test is primarily a test of $\mathbf{A}$.
6. **Gauge freedom.** Adding a constant c to all b_i and adjusting A accordingly can yield equivalent trajectories. The absolute scale of b vs. A is not independently identifiable from 16S data without growth-rate measurements.
7. **Joshi mapping approximation.** Filling 5 of 10 guilds from the Dieckow mean introduces a systematic bias toward the Dieckow community composition; the attractor analysis assumes stationarity (peri-implant communities near equilibrium) without direct verification.

5 Conclusions

We have shown that a metabolically informed three-layer sign prior significantly regularises Hamilton replicator dynamics for oral microbiome prediction. At the optimal AGORA weight $w = 1.0$, the model achieves:

- 100% sign agreement (35 constrained pairs, L1+L2+L3).
- LOO-RMSE = 0.0504 ± 0.021 (−14% vs. unconstrained gLV).
- LOO-BC = 0.147 (−5.3% vs. L1+L2, −47.7% vs. persistence null).
- Correct PCA/LDA community structure under prediction.

- CT-consistent patient-specific $\hat{b}$ clustering without supervision.
- Robust sign stability (SC= 1.00 for 31/45 pairs; 40/45 at SC $\geq$ 0.90) across all LOO folds.
- External validity: Dieckow **A** correctly orders peri-implant dysbiosis in 127 independent cross-sectional samples ($p < 0.0001$).

Key limitations are the small cohort ($N = 10$, $p = 0.08$ for the AGORA improvement), the symmetry constraint, and the caries-context prior. Future work should replicate in larger cohorts, evaluate asymmetric gLV with asymmetric sign priors, construct disease-specific priors for periodontitis, and incorporate posterior uncertainty via TMCMC [Nishioka et al., *in prep.*].

A AGORA2 representative strains and oral-fluid medium

Table 5: AGORA2 representative strains for each guild.

Guild	AGORA2 strain
Actinobacteria	<i>Actinomyces naeslundii</i> Howell 279
Bacilli	<i>Streptococcus gordonii</i> CH1
Negativicutes	<i>Veillonella parvula</i> Te3 DSM 2008
Bacteroidia	<i>Porphyromonas melaninogenica</i> ATCC 25845
Fusobacteriia	<i>Fusobacterium nucleatum</i> ATCC 25586
Clostridia	<i>Parvimonas micra</i> ATCC 33270
Coriobacteriia	<i>Atopobium parvulum</i> DSM 20469
Gammaproteobacteria	<i>Haemophilus parainfluenzae</i> T3T1
Betaproteobacteria	<i>Eikenella corrodens</i> ATCC 23834
Other	No AGORA representative (L3 absent)

Oral-fluid medium construction

The medium was assembled in two steps.

Step 1 — Literature-based saliva metabolite list. Quantified metabolite concentrations in unstimulated whole saliva were drawn from Dawes (2008) and Amerongen & Veerman (2002). Concentrations were converted from mM to mmol gDW⁻¹ h⁻¹ assuming a biofilm cell density of ~ 2 gDW L⁻¹ and a reference growth rate of 0.1 h⁻¹. The final list covers: sugars (glucose 10, fructose 5, sucrose 5, lactate 3 mmol gDW⁻¹ h⁻¹); 20 amino acids + Cys and Orn; B-vitamins (B₁, B₂, B₃, B₅, B₆ three forms, B₁₂, biotin, folate); trace metals (Fe^{2+/3+}, Mg²⁺, Ca²⁺, Zn²⁺, Cu²⁺, Mn²⁺, Co²⁺); purines (adenine, guanine); and inorganic salts (phosphate, sulfate, ammonium, Na⁺, K⁺, Cl⁻). Note: concentrations in v1 are set at the blood-plasma scale ($\sim 100\times$ realistic saliva values) to ensure all models remain feasible; a realistic v2 at actual salivary concentrations [Tenovuo, 1998] was tested but produced non-specific competition signals (81/90 guild pairs showed growth suppression $>50\%$) and was not used for the main analysis.

Step 2 — AGORA2-specific cofactors (verified by feasibility). Several AGORA2 genome-scale models require cofactors absent from saliva literature or present only at trace concentrations. These were identified iteratively: after adding the saliva-based components, each guild model was run with pFBA; models with objective value $< 10^{-6}$ were inspected for blocked reactions, and the missing exchange metabolite was added at a minimal bound. Components added in this step: protoheme & siroheme (required by *Porphyromonas*, *Prevotella*, obligate anaerobes); menaquinone MK-7/MK-8 (Gram-positive electron transport); meso-2,6-diaminopimelate (peptidoglycan precursor for Gram-negatives); saturated fatty acids C12/C14/C18 (membrane synthesis); 4-hydroxybenzoate (ubiquinone precursor); adenosylcobalamin; glutathione (oxidised/reduced,

required by Actinobacteria/Bacteroidia); putrescine, spermidine; cytidine (Gammaproteobacteria pyrimidine salvage). After both steps, all 9 guilds with an AGORA2 representative achieve $\mu \in [0.11, 1.66] \text{ h}^{-1}$ (verified). The exchange reaction identifiers use the AGORA2 BiGG legacy format (`EX_met(e)`).

B σ sensitivity of sign penalty

All models achieve identical sign agreement for $\sigma \in [0.05, 0.30]$, with LOO-RMSE variation $< 2\%$. We select $\sigma = 0.15$ as the operational value. At $\sigma = 0.05$ the prior approaches a hard constraint; at $\sigma = 0.50$ it is nearly uninformative. Results are qualitatively unchanged across the tested range.

C Parameter identifiability

155 parameters (55 for **A** + 100 for $\{\hat{\mathbf{b}}^{(p)}\}$) vs. 200 observations (ratio 1.29). Without regularisation the system is marginally overdetermined. The sign prior effectively removes one degree of freedom per constrained pair (35 pairs), raising the effective ratio to ≈ 1.67 . The LOO train-to-test RMSE gap is -0.007 to -0.014 for sign-prior models (prior helps generalisation) vs. $+0.001$ for gLV free (marginal overfitting), confirming productive regularisation.

D Null model Bray-Curtis comparison

Table 6: Null model LOO Bray-Curtis.

Model	LOO BC	vs. persistence
Persistence ($\hat{\phi}_t = \phi_{t-1}$)	0.281	—
Grand mean	0.254	−10%
Uniform ($1/K$)	0.617	+120%
gLV free	0.154	−45%
Hamilton L1+L2	0.155	−45%
Hamilton AGORA W1.0	0.147	−47.7%

LOO BC = 0.147 places predictions in the “same community type” range (< 0.2) relative to observations. The 11-point gap vs. L1+L2 (0.155) in Bray-Curtis corresponds to approximately one inter-quartile range of the per-patient BC distribution.

E 10-week MAP extrapolation

References

- Bashan A, Gibson TE, Friedman J, et al. (2016). Universality of human microbial dynamics. *Nature* 534:259–262.
- Bradbury J, Frostig R, Hawkins P, et al. (2018). JAX: composable transformations of Python+NumPy programs. <http://github.com/jax-ml/jax>
- Byrd RH, Lu P, Nocedal J, Zhu C (1995). A limited memory algorithm for bound constrained optimization. *SIAM Journal on Scientific Computing* 16(5):1190–1208.
- Cao H-T, Gibson TE, Bashan A, Liu Y-Y (2017). Inferring human microbial dynamics from temporal metagenomics data: pitfalls and lessons. *BioEssays* 39:1600188.

- Dawes C (2008). Salivary flow patterns and the health of hard and soft oral tissues. *Journal of the American Dental Association* 139 Suppl:18S–24S.
- Dieckow S, Szafranski SP, et al. (2024). Longitudinal dynamics of subgingival biofilm microbiome during caries recovery. *npj Biofilms and Microbiomes* 10:85.
- Fisher CK, Mehta P (2014). Identifying keystone species in the human gut microbiome from metagenomic timeseries using sparse linear regression. *PLoS ONE* 9:e102451.
- Heinken A, et al. (2023). Genome-scale metabolic reconstruction of 7,302 human microorganisms for personalized medicine. *Nature Biotechnology* 41:1320–1331.
- Hofbauer J, Sigmund K (1998). *Evolutionary Games and Population Dynamics*. Cambridge University Press.
- Joseph TA, et al. (2020). Compositional Lotka-Volterra describes microbial dynamics in health and disease. *PLOS Computational Biology* 16:e1007917.
- Joshi V, et al. (2025). Community-level profiling of peri-implant microbiome across health, mucositis and peri-implantitis. *mSystems* [in press].
- Junker P, Balzani D (2021). An extended Hamilton principle as unifying theory for coupled problems and dissipative microstructure evolution. *Continuum Mechanics and Thermodynamics* 33:1931–1956.
- Klempt F, Nishioka K, Soleimani M, Junker P (2024). Modelling biofilm growth and mechanics with the extended Hamilton principle. *Biomechanics and Modeling in Mechanobiology* 23.
- Klempt F, Nishioka K, Junker P (2025). Continuum biofilm model with replicator species dynamics. *[preprint]*
- Kolenbrander PE (2000). Oral microbial communities: biofilms, interactions, and genetic systems. *Annual Review of Microbiology* 54:413–437.
- Kolenbrander PE, Palmer RJ, Periasamy S, Jakubovics NS (2010). Oral multispecies biofilm development and the key role of cell-cell distance. *Nature Reviews Microbiology* 8:471–480.
- Kuntal BK, et al. (2019). ‘NetShift’: understanding ‘driver microbes’ from healthy and disease microbiome datasets. *ISME Journal* 13:442–454.
- Lewis NE, et al. (2010). Omic data from evolved *E. coli* are consistent with computed optimal growth from genome-scale models. *Molecular Systems Biology* 6:390.
- Mikx FHM, van der Hoeven JS (1975). Symbiosis of *Streptococcus mutans* and *Veillonella alcalescens* in mixed continuous cultures. *Archives of Oral Biology* 20:407–410.
- Nishioka K, Klempt F, Soleimani M, Junker P (in preparation). Bayesian inference for biofilm mechanics via TMCMC.
- Stein RR, et al. (2013). Ecological modelling from time-series inference: insight into dynamics and stability of intestinal microbiota. *PLOS Computational Biology* 9:e1003388.
- Szafranski SP, Nishioka K, et al. (2025). Metabolic interaction network for the peri-implant oral microbiome. *[preprint / in preparation]*
- Taylor PD, Jonker LB (1978). Evolutionarily stable strategies and game dynamics. *Mathematical Biosciences* 40:145–156.

- Chen T, Yu WH, Izard J, Baranova OV, Lakshmanan A, Dewhirst FE (2010). The Human Oral Microbiome Database: a web accessible resource for investigating oral microbe taxonomic and genomic information. *Database* 2010:baq013.
- Tenovuo, J. (1998). Antimicrobial function of human saliva — how important is it for oral health? *Acta Odontologica Scandinavica*, 56(5):250–256.
- Diener, C., Gibbons, S. M., and Resendis-Antonio, O. (2020). MICOM: Metagenome-Scale Modeling To Infer Metabolic Interactions in the Gut Microbiota. *mSystems*, 5(1):e00606-19.

Interaction matrix A — Hamilton + AGORA $W=1.0$
Training RMSE = 0.0565 | Pearson $r = 0.951$ | SA = 70/70




Figure 3: Fitted interaction matrix A (Hamilton replicator, AGORA $W = 1.0$, all 10 patients). Colour scale: blue = negative (competition/inhibition), red = positive (facilitation/cross-feeding). Diagonal $A_{ii} < 0$ (self-limitation) for all guilds. Green checkmarks ($\checkmark$): sign agrees with metabolic prior; red crosses ($\times$): discordant. Sign agreement = 70/70 (100%) for the full L1+L2+L3 model. Dominant positive interactions (red, top-left block): Actinobacteria–Bacilli and Bacilli–Betaproteobacteria reflect co-aggregation syntrophies [Kolenbrander, 2000]. Bacilli–Negativicutes is positive but muted ($A = +0.003$), consistent with the canonical *Streptococcus*–lactate–*Veillonella* cross-feeding axis [Mikx and van der Hoeven, 1975] operating on timescales longer than the 3-week sampling window.




Figure 6: Training trajectory prediction for a representative patient (Hamilton + AGORA W= 1.0). Solid lines: observed guild fractions; dashed: predicted. Overall Pearson $r = 0.874$ (all patients, all timepoints). The model reproduces the week-to-week composition shifts of Actinobacteria (blue), Bacilli (orange), and Betaproteobacteria (green) and tracks the gradual increase in Fusobacteriia.




Figure 7: Training fit for all 10 patients (Hamilton + AGORA W= 1.0). Solid bars: observed; hatched bars: predicted. Per-patient RMSE and Pearson r shown below each panel. Training RMSE = 0.057, $r = 0.951$, SA = 70/70 (100%). Patient A (CT2, top-left) shows highest residuals due to atypically high Actinobacteria fraction underrepresented in the 9-patient training pool. Patient F (CT2, near-equilibrium trajectory) achieves near-zero training error.




Figure 8: LOO-RMSE per patient for L1+L2 (grey) vs. AGORA W= 1.0 (blue). Mean LOO-RMSE: L1+L2 = 0.0516, AGORA W= 1.0 = 0.0504 (-2.4%). 7/10 patients improve; Patient L is the only clear regression. CT (CT1/CT2) and RMSE values annotated. Horizontal dashed line: gLV free baseline (0.0588).




Figure 9: LOO Bray-Curtis dissimilarity for all model variants and null models. Null baselines (dashed): persistence (0.281), grand-mean (0.254). All ODE models substantially beat both nulls. Hamilton AGORA W= 1.0 (blue) achieves the lowest BC (0.147), 5.3% below L1+L2 and 47.7% below persistence. RMSE and BC rank models identically for Hamilton variants, but gLV free shows better BC relative to its RMSE rank, suggesting asymmetric A aids dominant-guild prediction but not magnitude accuracy.




Figure 10: RMSE vs. Bray-Curtis scatter across all model variants and LOO folds. Each point is one patient-fold. Axes: LOO-RMSE (x) and LOO-BC (y). Model colour as in legend. Patients with large RMSE typically also have large BC (Pearson $r = 0.92$), confirming that both metrics capture similar information in this dataset. Hamilton AGORA $W = 1.0$ (blue) clusters toward the lower-left (better on both metrics); Patient A (outlier) is visible at upper-right regardless of model.

MICOM Tradeoff Fraction τ Sensitivity (gLV model)




Figure 11: MICOM tradeoff fraction τ sensitivity (gLV model, 10-fold LOO-CV). **Left:** mean LOO-RMSE vs. $\tau \in \{0.3, 0.4, 0.5\}$. Dashed lines indicate gLV baselines (green: $\alpha=0.25$, blue: AGORA v1). **Right:** sign agreement (of 72 constrained pairs). LOO-RMSE varies by <0.001 across τ ; $\tau = 0.5$ (Diener et al. 2020 default) is adopted.




Figure 12: Per-patient LOO-RMSE for Hamilton model variants: L1+L2 (grey), AGORA v1 (blue), and MICOM $\tau=0.5$ (red). Dashed horizontal lines show respective means. MICOM achieves the lowest mean RMSE (0.0502) and improves over AGORA v1 in 6/10 patients.




Figure 13: Ordination of observed and LOO-predicted W2/W3 compositions (Hamilton + AGORA W= 1.0, 10 patients). **(A)** PCA: predicted points (open symbols) shadow their corresponding observed points (solid), confirming that the model reproduces the compositional data manifold and does not project to a degenerate region. **(B)** LDA with CT1/CT2 labels: the linear discriminant separating community types is preserved under prediction (LDA accuracy: observed = 100%; predicted = 90%, patient A misclassified due to Actinobacteria overestimation). Coloured by patient; CT1 = triangles, CT2 = circles.




Figure 14: PCA of patient-specific growth vectors $\hat{b}^{(p)}$ (Hamilton + AGORA W= 1.0). PC1 (explained variance shown) separates CT1 (triangles) from CT2 (circles) without label supervision, demonstrating that patient-level ecological differences are primarily captured in the intrinsic growth rates, not in the shared **A**. Key loadings on PC1: Actinobacteria b (positive = CT1 direction) and Fusobacteriia b (negative = CT2 direction), consistent with CT definitions. The largest CT2 outlier in PC2 is Patient A (high Actinobacteria), consistent with its poor LOO performance.




Figure 15: A matrix LOO stability (gLV $\alpha=0.25$, best model, 10 folds). **(A)** Sign consistency (SC) heatmap; boxed cells have a non-zero sign prior. 31/45 pairs: SC = 1.00; 40/45: SC ≥ 0.90 ; 1 pair: SC < 0.70. **(B)** $|\bar{A}_{ij}|$ vs. LOO std, coloured by regime (blue=data+prior aligned, light blue=prior-muted, red=data-driven). **(C)** Fold-to-fold distributions for the 12 strongest pairs.





Figure 16: AGORA weight sweep ($w_{L3} \in \{0, 0.5, 1.0, 1.5, 2.0\}$). Panels: **(A)** training RMSE (mean $\pm$ std, LOO folds); **(B)** Pearson r ; **(C)** sign agreement (SA). A discrete phase transition at $w = 1.0$: SA jumps from 94% to 100% with a simultaneous drop in training RMSE (0.063 $\rightarrow$ 0.057). At $w > 1.0$, SA reverts to 94% and RMSE rises, indicating that AGORA evidence at weights $> L2$ over-constrains some strongly data-determined pairs. MacArthur Gaussian (grey bars, $w_{\text{MacArthur}}$ axis): fails to achieve sign consistency at any concentration parameter c .

External validation: peri-implant dysbiosis prediction
Dieckow A matrix $\rightarrow$ Joshi et al. 2025 (n=127)



AGORA W=1.0 Hamilton ODE: 10-week extrapolation
Dots = observed (wk1-3); lines = predicted; shaded = extrapolation region




Figure 18: MAP trajectory extrapolation to week 10 (AGORA W= 1.0, all 10 patients). Training window (weeks 1–3) shaded; dots = observed. Solid lines = forward integration with per-patient $\hat{b}$. CT1 patients converge to Actinobacteria/Bacilli-dominated equilibria; CT2 patients to Fusobacteriia/Bacteroidia enrichment. Convergence by weeks 5–7 in most patients. This extrapolation is speculative (no longitudinal data beyond week 3 available for validation) and should be interpreted as a qualitative attractor projection, not a clinical prediction.